

Hw4_p3

Qiuying Li UNI ql2280

3/28/2017

```
setwd("~/Desktop/2017 spring/GR 5241/HW/HW4")
library("freestats")
# Part 1
#step1
#We need to calculate data points which are currently classified correctly and are weighted down
#First of all, we are giving the weights = 1/n
#yi = {-1, +1} are class labels.
#stump ignores all entries of x except xj, it is equivalent to a linear classifier defined by an affine hyperplane.
#The plane is orthogonal to the jth axis, with which it intersects at xj = theta.
#The orientation of the hyperplane is determined by m = {-1, +1}.

#In order to have the parameters j, m and theta, we need a decisionStump function
train_parameters = function (X,w,Y){
  parameters = decisionStump(X,w,Y)
  j = parameters$j
  theta = parameters$theta
  m = parameters$m
  return(j, theta, m)
}

# We want to generate a decisionStump function between -1 and 1
# w is the weight which is 1/n
# The goal of this function is to calculate optimal theta, m and j in all dimensions.
train = function (X,w,Y){
  a = nrow(X)
  b = ncol(X)
  c_min= rep(a,b)
  theta_min = rep(-2,b)
  m_min = rep(-1,p)
  #In order to have the optimal theta, we need to build cost functions for -1 and 1
  #when yi ∈ { -1, 1}
  cost_1 = function (theta,x,y,number,weights){
    classifier = red(-1, number)
    classifier[x > theta] = 1
    result_sum = sum(weights*classifier != y)
    return(result_sum)
  }
```

```

}

#when  $y_i \in \{1\}$ 
cost_neg1 = function (theta,x,y,number,weights){
  classifier = red(1, number)
  classifier[x < theta] = -1
  result_sum = sum(weights*classify != y)
  return(result_sum)
}

# compute optimal theta for each dimension
for (dim in seq(b)) {
  X_dim <- X[,dim]
  unique_X_dim <- unique(X_dim)
  unique_X_dim <- c(unique_X_dim, -2)
  # costs for both  $m = 1$  and  $-1$ 
  cost_m_1 <- apply(matrix(unique_X_dim), 1, cost_1, x=X_dim, y=Y, number=n
, weights=w)
  cost_m_neg1 <- apply(matrix(unique_X_dim), 1, cost_neg1, x=X_dim, y=Y, nu
mber=n, weights=w)

  if (min(cost_m_neg1) < min(cost_m_1)) {
    ind <- which.min(cost_m_neg1)
    m_min[dim] <- -1
    c_min[dim] <- c_d_n[ind]
  } else {
    ind <- which.min(cost_m_1)
    c_min[d] <- cost_m_1[ind]
  }
  theta_min[d] <- unique_X_dim[ind]
}

# find out the dimension with the optimal theta and m
min_dim <- which.min(c_min)
theta_min <- theta_min[min_d]
c_min <- c_min[min_d]
m_min <- m_min[min_d]
return(c(min_dim, theta_min, m_min))
}

#Step 2
#label <- classify(X, pars) for the classification routine,
#which evaluates the weak learner on X using the parametrization pars.
Pars are the parameters of the <j,m,theta>

classify <- function(X, pars) {
  # classify X use the parameters in pars
  j <- pars[1]

```

```

theta <- pars[2]
m <- pars[3]
n <- nrow(X)
X_d <- X[,j]
classified <- rep(-m, n)
classified[X_d > theta] = m
return(matrix(classified))
}

#step 3
#A function c hat <- agg class(X, alpha, all_pars)
# It evaluates the boosting classifier ("ag- gregated classifier") on X.
#The argument alpha denotes the vector of voting weights and all_pars contain
s the parameters of all weak learners.

agg_class <- function(X, alpha, all_pars) {
  a<- nrow(X)
  b <- length(alpha)
  agg_labels <- matrix(0, a, 1)

  # we need to be careful when there is only one row
  # For this case, the function is no longer applicable.
  if (b == 1) {
    all_pars <- rbind(all_pars, matrix(0,1,3))
  }

  # otherwise, when we have multiple of rows
  for (i in seq(b)) {
    al <- alpha[i]
    agg_labels <- agg_labels + al * classify(X, all_pars[i,])
  }
  classified <- matrix(-1, a, 1)
  classified[agg_labels >= 0] <- 1
  return(classified)
}

# Part 2
# Implement the functions train and classify for decision stumps.
comb_fun <- function(X, Y, all_pars, alpha, iter) {
  # In order to perform a cross validation, we choose k = 5
  k = 5
  fold_size <- n / k
  cv_errors <- matrix(1,k,1)
  for (cv in seq(k)) {
    cv_j <- X[-(((cv-1)*fold_size+1):(cv*fold_size)),]
    cv_m <- Y[-(((cv-1)*fold_size+1):(cv*fold_size)),]
    cv_theta <- w[-(((cv-1)*fold_size+1):(cv*fold_size)),]

```

```

cv_set_j <- X[((cv-1)*fold_size+1):(cv*fold_size),]
cv_set_m <- Y[((cv-1)*fold_size+1):(cv*fold_size),]

# After cross validation, we can plug in the 3 parameters,m,j theta in train function and classify function

cv_pars <- train(cv_j, cv_theta, cv_m)
tr_pred_labels <- classify(cv_j, cv_pars)

# calculate the prediction error rate
tr_error_rate <- sum(cv_theta*(tr_pred_labels != tr_set_labels)) / sum(cv_theta)

# Compute voting weights alpha by given formula
cv_alpha <- log((1-tr_error_rate)/tr_error_rate)
cv_labels <- agg_class(cv_set_feats, c(alpha[0:(iter-1)],), cv_alpha), rbind(all_pars[0:(iter-1)],), cv_pars))

# compute cross validation error rate
cv_error <- sum(cv_labels != cv_set_labels) / fold_size
cv_errors[cv,] <- cv_error
}
cv_avg_error <- mean(cv_errors)
return(cv_avg_error)
}

# part 3
#step 1 read USPS data
train3 <- as.matrix(read.table("train_3.txt", header=FALSE, sep=","))
train8 <- as.matrix(read.table("train_8.txt", header=FALSE, sep=","))
xtrain <- rbind(train3, train8)
ytrain3 <- rep(c(1,-1), c(nrow(train3), nrow(train8)))
ytrain3 <- matrix(ytrain3)
test <- as.matrix(read.table("zip_test.txt",header = F))
test <- test[test[,1]%in%c(3,8),]
xtest <- test[,-1]
ytest <- test[,1]
ytest[ytest == 3] <- 1
ytest[ytest == 8] <- -1
ytest <- matrix(ytest)

#step 2: apply AdaBoost function to USPS data
# To calculate the test error, training error and cross validation error
AdaBoost <- function(B, X, Y, testX, testY){
  n <- nrow(X)
  test_size <- nrow(testX)
  w <- matrix(1/n, n)
  alpha <- matrix(0, B, 1)

```

```

all_pars <- matrix(0, B, 3)
error <- matrix(0, B, 3)
iterator <- 0
# we need to perform a 5-cross validation through combined train function
and classify function, to find out the three parameters
while (iterator < B) {
  iterator = iterator + 1
  cv_error <- comb_fun (X, Y, all_pars, alpha, iterator)
  # calculate the three parameters m,j,theta
  pars <- train(X, w, Y)
  # use above parameters to perform classify function
  labels <- classify(X, pars)
  # calculate error of the training data set
  error_rate <- sum(w*(labels != Y)) / sum(w)
  # compute voting weights alpha
  alpha <- log((1-error_rate)/error_rate)
  # plug the calculated parameters and voting weights
  alpha[iterator,] <- alpha
  all_pars[iterator,] <- pars
  # recompute weights w
  w <- w * exp(alpha * (labels != Y))
  # calculate error of the test data
  label_test <- agg_class(testX, alpha[1:iterator,], all_pars[1:iterator,])
  error_test <- sum(label_test != testY) / test_size
  error[iterator,1] <- error_rate
  error[iterator,2] <- cv_error
  error[iterator,3] <- error_test
}
return(cbind(alpha, all_pars,error))
}
n <- nrow(xtrain)
w <- matrix(1/n, n)

result <- AdaBoost(100, xtrain, ytrain3, xtest, ytest)

#part 4: make a plot of the error

test_error = result[,7]

training_error = result[,6]

plot(training_error,type = "l",lty = 1,ylim=c(0, 0.2),xlab = "Iterator", ylab
= "Error Rate"

lines(test_error, lty = 6)

legend("topright", c("training error" ,"testing error"), lty=c(1, 6),cex=0.7)

```

```
mtext(1, text="Iteration", cex=0.6, line = 1)
```



