

Lecture 14: Decision trees





Reading: Section 9.2

GU4241/GR5241 Statistical Machine Learning

Linxi Liu


March 2, 2017

Decision trees, 10,000 foot view



1. Find a partition of the space of predictors.
2. Predict a constant in each set of the partition.
3. The partition is defined by splitting the range of one predictor at a time.
→ Not all partitions are possible.

Example: Predicting a baseball player's salary



The prediction for a point in R_i is the average of the training points in R_i .


How is a decision tree built?

- ▶ Start with a single region R_1 , and iterate:
 1. Select a region R_k , a predictor X_j , and a splitting point s , such that splitting R_k with the criterion $X_j < s$ produces the largest decrease in RSS:

$$\sum_{m=1}^{|T|} \sum_{x_i \in R_m} (y_i - \bar{y}_{R_m})^2$$

2. Redefine the regions with this additional split.
- ▶ Terminate when there are 5 observations or fewer in each region.
 - ▶ This grows the tree from the root towards the leaves.

How is a decision tree built?



How do we control overfitting?

- ▶ **Idea 1:** Find the optimal subtree by cross validation.
 - There are too many possibilities, so we would still over fit.
- ▶ **Idea 2:** Stop growing the tree when the RSS doesn't drop by more than a threshold with any new cut.
 - In our greedy algorithm, it is possible to find good cuts after bad ones.

How do we control overfitting?

- ▶ **Cost complexity pruning:**

- ▶ Solve the problem:

$$\text{minimize } \sum_{m=1}^{|T|} \sum_{x_i \in R_m} (y_i - \bar{y}_{R_m})^2 + \alpha |T|.$$



- ▶ When $\alpha = \infty$, we select the null tree.
 - ▶ When $\alpha = 0$, we select the full tree.
 - ▶ The solution for each α is among T_1, T_2, \dots, T_m from weakest link pruning.
 - ▶ Choose the optimal α (the optimal T_i) by cross validation.

Cross validation



1. Split the training points into 10 folds.
2. For $k = 1, \dots, 10$, using every fold except the k th:
 - ▶ Construct a sequence of trees T_1, \dots, T_m for a range of values of α , and find the prediction for each region in each one.
 - ▶ For each tree T_i , calculate the RSS on the test set.
3. Select the parameter α that minimizes the average test error.

Note: We are doing all fitting, **including the construction of the trees**, using only the training data.

Example. Predicting baseball salaries



Example. Predicting baseball salaries



Classification trees

- ▶ They work much like regression trees.
- ▶ We predict the response by **majority vote**, i.e. pick the most common class in every region.
- ▶ Instead of trying to minimize the RSS:

$$\sum_{m=1}^{|T|} \sum_{x_i \in R_m} (y_i - \bar{y}_{R_m})^2$$

we minimize a classification loss function.

Classification losses

- ▶ The 0-1 loss or misclassification rate:

$$\sum_{m=1}^{|T|} \sum_{x_i \in R_m} \mathbf{1}(y_i \neq \hat{y}_{R_m})$$

- ▶ The Gini index:

$$\sum_{m=1}^{|T|} q_m \sum_{k=1}^K \hat{p}_{mk}(1 - \hat{p}_{mk}),$$

where $\hat{p}_{m,k}$ is the proportion of class k within R_m , and q_m is the proportion of samples in R_m .

- ▶ The cross-entropy:

$$- \sum_{m=1}^{|T|} q_m \sum_{k=1}^K \hat{p}_{mk} \log(\hat{p}_{mk}).$$

Classification losses







Figure: Node impurity measures for two-class classification

Classification losses

- ▶ The Gini index and cross-entropy are better measures of the purity of a region, i.e. they are low when the region is mostly one category.
- ▶ **Motivation for the Gini index:**

If instead of predicting the most likely class, we predict a random sample from the distribution $(\hat{p}_{1,m}, \hat{p}_{2,m}, \dots, \hat{p}_{K,m})$, the Gini index is the expected misclassification rate.
- ▶ It is typical to use the Gini index or cross-entropy for growing the tree, while using the misclassification rate when pruning the tree.

Example. Heart dataset.



Some advantages of decision trees


- ▶ Very easy to interpret!
- ▶ Closer to human decision-making.
- ▶ Easy to visualize graphically.
- ▶ They easily handle qualitative predictors and missing data.

Classification and Regression trees, in a nut shell

- ▶ Grow the tree by recursively splitting the samples in the leaf R_i according to $X_j > s$, such that (R_i, X_j, s) maximize the drop in RSS.
→ Greedy algorithm.
- ▶ Create a sequence of subtrees T_0, T_1, \dots, T_m using a **pruning** algorithm.
- ▶ Select the best tree T_i (or the best α) by cross validation.
→ Why might it be better to choose α instead of the tree T_i by cross-validation?

Example. Heart dataset.

How do we deal with categorical predictors?



Categorical predictors

- ▶ If there are only 2 categories, then the split is obvious. We don't have to choose the splitting point s , as for a numerical variable.
- ▶ If there are more than 2 categories:
 - ▶ Order the categories according to the average of the response:

ChestPain : a > ChestPain : c > ChestPain : b

- ▶ Treat as a numerical variable with this ordering, and choose a splitting point s .
- ▶ One can show that this is the optimal way of partitioning.

Missing data

- ▶ Suppose we can assign every sample to a leaf R_i despite the missing data.
- ▶ When choosing a new split with variable X_j (growing the tree):
 - ▶ Only consider the samples which have the variable X_j .
 - ▶ In addition to choosing the best split, choose a second best split using a different variable, and a third best, ...
- ▶ To propagate a sample down the tree, if it is missing a variable to make a decision, try the second best decision, or the third best, ...