

Lecture 15: Bagging, Random Forests

Reading: Sections 9.2, 15.2, 15.3

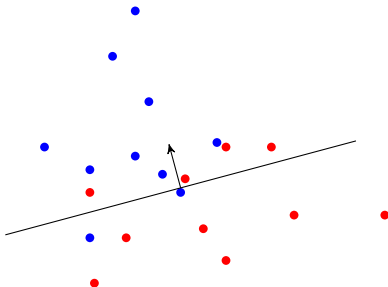
GU4241/GR5241 Statistical Machine Learning

Linxi Liu

March 7, 2017

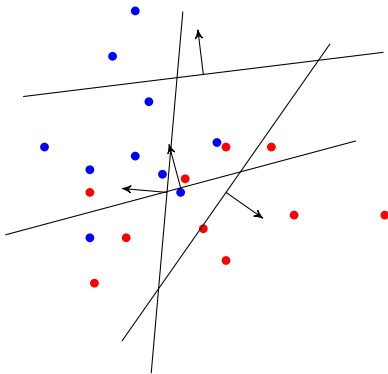
Ensembles

A *randomly* chosen hyperplane classifier has an *expected* error of 0.5 (i.e. 50%).



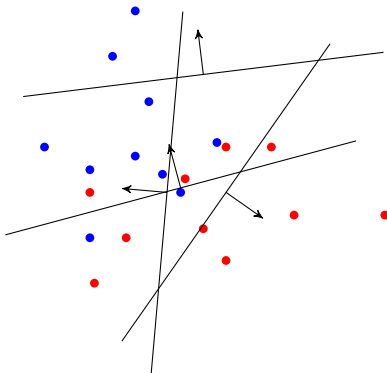
Ensembles

A *randomly* chosen hyperplane classifier has an *expected* error of 0.5 (i.e. 50%).



Ensembles

A *randomly* chosen hyperplane classifier has an *expected* error of 0.5 (i.e. 50%).



- ▶ Many random hyperplanes combined by majority vote: Still 0.5.
- ▶ A single classifier slightly better than random: $0.5 + \epsilon$.
- ▶ What if we use m such classifiers and take a majority vote?

Voting

Decision by majority vote

- ▶ m individuals (or classifiers) take a vote. m is an odd number.
- ▶ They decide between two choices; one is correct, one is wrong.
- ▶ After everyone has voted, a decision is made by simple majority.

Note: For two-class classifiers f_1, \dots, f_m (with output ± 1):

$$\text{majority vote} = \text{sgn}\left(\sum_{j=1}^m f_j\right)$$

Assumptions

Before we discuss ensembles, we try to convince ourselves that voting can be beneficial. We make some simplifying assumptions:

- ▶ Each individual makes the right choice with probability $p \in [0, 1]$.
- ▶ The votes are *independent*, i.e. stochastically independent when regarded as random outcomes.

Does the Majority Make the Right Choice?

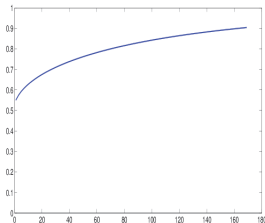
Condorcet's rule

If the individual votes are independent, the answer is

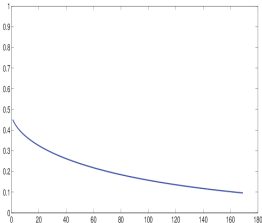
$$\Pr\{\text{majority makes correct decision}\} = \sum_{j=\frac{m+1}{2}}^m \frac{m!}{j!(m-j)!} p^j (1-p)^{m-j}$$

This formula is known as **Condorcet's jury theorem**.

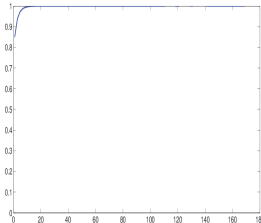
Probability as function of the number of votes



$p = 0.55$



$p = 0.45$



$p = 0.85$

Ensemble Methods

Terminology

- ▶ An **ensemble method** makes a prediction by combining the predictions of many classifiers into a single vote.
- ▶ The individual classifiers are usually required to perform only slightly better than random. For two classes, this means slightly more than 50% of the data are classified correctly. Such a classifier is called a **weak learner**.

Strategy

- ▶ We have seen above that if the weak learners are random and independent, the prediction accuracy of the majority vote will increase with the number of weak learners.
- ▶ Since the weak learners all have to be trained on the training data, producing random, independent weak learners is difficult.
- ▶ Different ensemble methods (e.g. Boosting, Bagging, etc) use different strategies to train and combine weak learners that behave relatively independently.

Methods We Will Discuss

Boosting

- ▶ After training each weak learner, data is modified using weights.
- ▶ Deterministic algorithm.

Bagging

Each weak learner is trained on a random subset of the data.

Random forests

- ▶ Bagging with decision trees as weak learners.
- ▶ Uses an additional step to remove dimensions in \mathbb{R}^d that carry little information.

Bagging = Bootstrap Aggregation

- ▶ Replicate the dataset by sampling with replacement.
- ▶ We apply a learning method to each bootstrap replicate, to produce predictions $\hat{f}^{(1)}, \dots, \hat{f}^{(B)}$.
- ▶ In Chapter 5, we were interested in the variability of these predictions:

$$\text{SE}(\hat{f}(x)) \approx \text{SD}(\hat{f}^{(1)}(x), \dots, \hat{f}^{(B)}(x)).$$

- ▶ Now, we will use the average of these predictions as an estimator with reduced variance:

$$\hat{f}^{\text{bag}}(x) = \frac{1}{B} \sum_{b=1}^B \hat{f}^{(b)}(x)$$

we treat all x as independent, only average is enough

Bagging decision trees

trees suffer from the high variance

- ▶ Replicate the dataset by sampling with replacement.
- ▶ Fit a decision tree to each bootstrap replicate (growing the tree, and pruning).
- ▶ **Regression:** To make a prediction for an input point x , average the predictions of all the trees:

$$\hat{f}^{\text{bag}}(x) = \frac{1}{B} \sum_{b=1}^B \hat{f}^{(b)}(x)$$

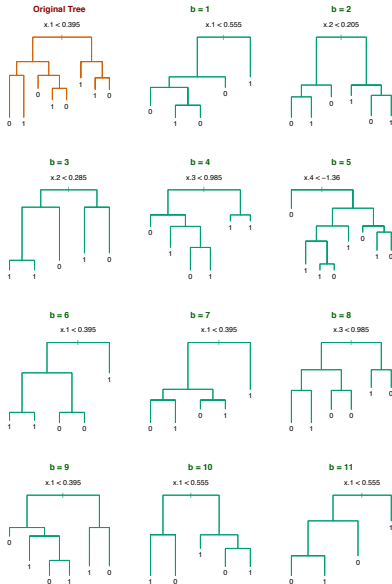
- ▶ **Classification:** To make a prediction for an input point x_0 , take the majority vote from the set of predictions:

$$\hat{y}_0^{(1)}, \dots, \hat{y}_0^{(B)}.$$

Example: Bagging decision trees

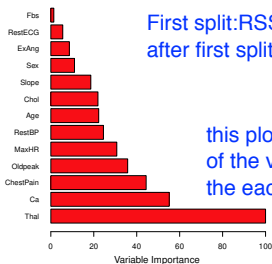
each tree corresponds to
each bootstrap

- ▶ Two classes, each with Gaussian distribution in \mathbb{R}^5 .
- ▶ Note the variance between bootstrapped trees.



Bagging decision trees

- ▶ **Disadvantage:** Every time we fit a decision tree to a Bootstrap sample, we get a different tree T^b .
 - **Loss of interpretability** it is hard to quantify which predictors is important, which is we usually want
- ▶ For each predictor, add up the total amount by which the RSS (or Gini index) decreases every time we use the predictor in T^b .
- ▶ Average this total over each Bootstrap estimate T^1, \dots, T^B .



First split: $RSS_0 - RSS_{1L} - RSS_{0r}$
after first split: $RSS_{0r} - RSS_{1r}$

this plot; account for the amount
of the variation of
the each predictors

disadvantages of bagging: **Out-of-bag (OOB) error**
trees can be very similar to each other for bagging.
means trees are highly correlated, then the performance
might only focus on one respect.

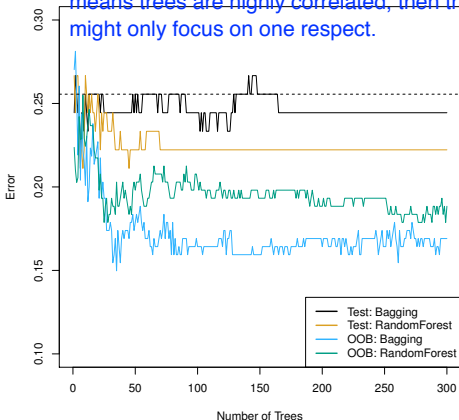
- ▶ To estimate the test error of a bagging estimate, we could use cross-validation.
- ▶ Each time we draw a Bootstrap sample, we only use 63% of the observations.
- ▶ **Idea:** use the rest of the observations as a test set.
- ▶ **OOB error:**
 - ▶ For each sample x_i , find the prediction \hat{y}_i^b for all bootstrap samples b which do not contain x_i . There should be around $0.37B$ of them. Average these predictions to obtain \hat{y}_i^{oob} .
 - ▶ Compute the error $(y_i - \hat{y}_i^{\text{oob}})^2$.
 - ▶ Average the errors over all observations $i = 1, \dots, n$.
- ▶ For B large, OOB error is virtually equivalent to LOOCV.

Out-of-bag (OOB) error

disadvantages of bagging:

trees can be very similar to each other for bagging.

means trees are highly correlated, then the performance might only focus on one respect.



The test error decreases as we increase B
(dashed line is the error for a plain decision tree).

Random Forests

Bagging has a problem:

→ The trees produced by different Bootstrap samples can be very similar.

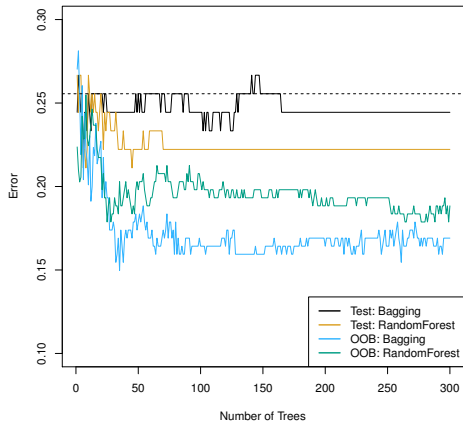
if we have p predictors in total, each time we draw 3 samples, and construct a tree based on it;
next, draw different 3 samples and build trees

Random Forests:

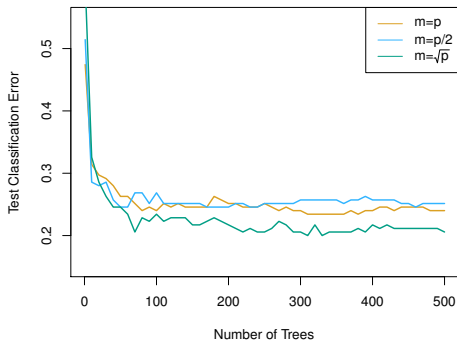
- ▶ We fit a decision tree to different Bootstrap samples.
- ▶ When growing the tree, we select a random sample of $m < p$ predictors to consider in each step.
- ▶ This will lead to very different (or “uncorrelated”) trees from each sample.
- ▶ Finally, average the prediction of each tree.

usually, random forest is better than bagging

Random Forests vs. Bagging



Random Forests, choosing m



The optimal m is usually around \sqrt{p} ,
but this can be used as a tuning parameter.